



Vidyavardhini's College of Engineering &
Technology

Department of Computer Engineering

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 07-08-2023
Date of Submission: 08-10-2023



Vidyavardhini's College of Engineering & Technology

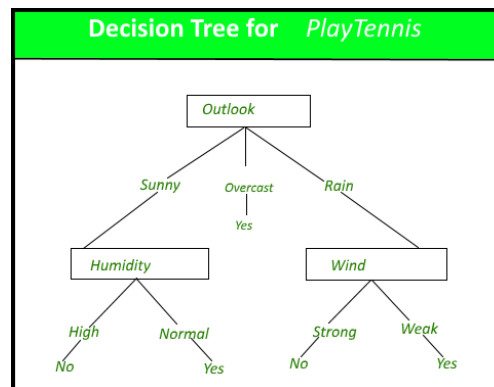
Department of Computer Engineering

Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

ml-experiment-3

October 8, 2023

```
[ ]: # Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # Adult dataset path
adult_dataset_path = "adult.csv"

# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)
```

```
[ ]: # Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

```
[ ]:  age workclass  fnlwgt      education  education.num marital.status \
0   90      ?      77053      HS-grad              9      Widowed
1   82  Private  132870      HS-grad              9      Widowed
2   66      ?  186061  Some-college             10      Widowed

      occupation  relationship  race  sex  capital.gain  capital.loss \
0              ?  Not-in-family  White  Female          0         4356
1  Exec-managerial  Not-in-family  White  Female          0         4356
2              ?      Unmarried  Black  Female          0         4356

hours.per.week  native.country  income
0              40  United-States  <=50K
```

```

1          18 United-States <=50K
2          40 United-States <=50K

```

```

[ ]: print ("Rows      : " ,df.shape[0])
      print ("Columns   : " ,df.shape[1])
      print ("\nFeatures : \n" ,df.columns.tolist())
      print ("\nMissing values : ", df.isnull().sum().values.sum())
      print ("\nUnique values : \n",df.nunique())

```

```

Rows      : 32561
Columns   : 15

```

```

Features :
['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status',
'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss',
'hours.per.week', 'native.country', 'income']

```

```

Missing values : 0

```

```

Unique values :
age          73
workclass     9
fnlwgt       21648
education     16
education.num 16
marital.status 7
occupation    15
relationship   6
race          5
sex           2
capital.gain  119
capital.loss   92
hours.per.week 94
native.country 42
income        2
dtype: int64

```

```

[ ]: # Let's understand the type of values present in each column of our adult_
      ↪dataframe 'df'.
      df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             32561 non-null  int64

```

```

1  workclass      32561 non-null  object
2  fnlwgt        32561 non-null  int64
3  education     32561 non-null  object
4  education.num  32561 non-null  int64
5  marital.status 32561 non-null  object
6  occupation    32561 non-null  object
7  relationship  32561 non-null  object
8  race          32561 non-null  object
9  sex           32561 non-null  object
10 capital.gain  32561 non-null  int64
11 capital.loss  32561 non-null  int64
12 hours.per.week 32561 non-null  int64
13 native.country 32561 non-null  object
14 income       32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
[ ]: # Numerical feature of summary/description
df.describe()
```

```
[ ]:
count    32561.000000  3.256100e+04  32561.000000  32561.000000  32561.000000  \
mean      38.581647  1.897784e+05      10.080679    1077.648844     87.303830
std       13.640433  1.055500e+05      2.572720    7385.292085    402.960219
min       17.000000  1.228500e+04      1.000000      0.000000      0.000000
25%       28.000000  1.178270e+05      9.000000      0.000000      0.000000
50%       37.000000  1.783560e+05     10.000000      0.000000      0.000000
75%       48.000000  2.370510e+05     12.000000      0.000000      0.000000
max       90.000000  1.484705e+06     16.000000   99999.000000   4356.000000

      hours.per.week
count    32561.000000
mean      40.437456
std       12.347429
min        1.000000
25%       40.000000
50%       40.000000
75%       45.000000
max       99.000000

```

```
[ ]: # pull top 5 row values to understand the data and how it's look like
df.head()
```

```
[ ]:
   age workclass  fnlwgt  education  education.num marital.status  \
0   90      ?      77053      HS-grad           9      Widowed
1   82  Private  132870      HS-grad           9      Widowed
2   66      ?  186061  Some-college          10      Widowed

```

3	54	Private	140359	7th-8th	4	Divorced
4	41	Private	264663	Some-college	10	Separated

	occupation	relationship	race	sex	capital.gain	\
0	?	Not-in-family	White	Female	0	
1	Exec-managerial	Not-in-family	White	Female	0	
2	?	Unmarried	Black	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	

	capital.loss	hours.per.week	native.country	income
0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K

```
[ ]: # checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass
```

```
[ ]: 1836
```

```
[ ]: # checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()
df_check_missing_occupation
```

```
[ ]: 1843
```

```
[ ]: # checking "?" values, how many are there in the whole dataset
df_missing = (df=='?').sum()
df_missing
```

```
[ ]: age                0
workclass             1836
fnlwgt                0
education             0
education.num         0
marital.status        0
occupation            1843
relationship          0
race                  0
sex                   0
capital.gain          0
capital.loss          0
hours.per.week        0
native.country        583
```

```
income          0
dtype: int64
```

```
[ ]: percent_missing = (df=='?').sum() * 100/len(df)
percent_missing
```

```
[ ]: age          0.000000
workclass        5.638647
fnlwgt           0.000000
education        0.000000
education.num    0.000000
marital.status   0.000000
occupation       5.660146
relationship     0.000000
race             0.000000
sex              0.000000
capital.gain     0.000000
capital.loss     0.000000
hours.per.week   0.000000
native.country   1.790486
income           0.000000
dtype: float64
```

```
[ ]: #Let's find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()
```

```
[ ]: age          32561
workclass        30725
fnlwgt           32561
education        32561
education.num    32561
marital.status   32561
occupation       30718
relationship     32561
race             32561
sex              32561
capital.gain     32561
capital.loss     32561
hours.per.week   32561
native.country   31978
income           32561
dtype: int64
```

```
[ ]: # dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()
```



```
[ ]:  age workclass  fnlwgt      education  education.num marital.status \
1    82   Private  132870      HS-grad           9      Widowed
3    54   Private  140359      7th-8th           4      Divorced
4    41   Private  264663  Some-college          10      Separated
5    34   Private  216864      HS-grad           9      Divorced
6    38   Private  150601      10th             6      Separated

      occupation  relationship  race    sex  capital.gain \
1   Exec-managerial  Not-in-family  White  Female           0
3  Machine-op-inspct    Unmarried  White  Female           0
4   Prof-specialty    Own-child  White  Female           0
5   Other-service    Unmarried  White  Female           0
6   Adm-clerical    Unmarried  White   Male           0

      capital.loss  hours.per.week  native.country  income
1           4356           18  United-States  <=50K
3           3900           40  United-States  <=50K
4           3900           40  United-States  <=50K
5           3770           45  United-States  <=50K
6           3770           40  United-States  <=50K
```

```
[ ]: # select all categorical variables
df_categorical = df.select_dtypes(include=['object'])

# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()
```

```
[ ]: workclass      0
     education      0
     marital.status  0
     occupation      7
     relationship    0
     race            0
     sex             0
     native.country  556
     income          0
     dtype: int64
```

```
[ ]: # dropping the "?"s from occupation and native.country
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	age	30162 non-null	int64
1	workclass	30162 non-null	object
2	fnlwgt	30162 non-null	int64
3	education	30162 non-null	object
4	education.num	30162 non-null	int64
5	marital.status	30162 non-null	object
6	occupation	30162 non-null	object
7	relationship	30162 non-null	object
8	race	30162 non-null	object
9	sex	30162 non-null	object
10	capital.gain	30162 non-null	int64
11	capital.loss	30162 non-null	int64
12	hours.per.week	30162 non-null	int64
13	native.country	30162 non-null	object
14	income	30162 non-null	object

dtypes: int64(6), object(9)

memory usage: 3.7+ MB

```
[ ]: from sklearn import preprocessing

# encode categorical variables using label Encoder

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

```
[ ]: workclass    education marital.status    occupation  relationship \
1    Private      HS-grad      Widowed    Exec-managerial Not-in-family
3    Private      7th-8th      Divorced    Machine-op-inspct  Unmarried
4    Private    Some-college    Separated    Prof-specialty    Own-child
5    Private      HS-grad      Divorced    Other-service    Unmarried
6    Private      10th        Separated    Adm-clerical     Unmarried
```

```
      race    sex native.country  income
1  White  Female  United-States  <=50K
3  White  Female  United-States  <=50K
4  White  Female  United-States  <=50K
5  White  Female  United-States  <=50K
6  White   Male   United-States  <=50K
```

```
[ ]: # apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

```
[ ]:  workclass  education  marital.status  occupation  relationship  race  sex  \
1         2         11           6           3           1      4    0
3         2         5           0           6           4      4    0
4         2        15           5           9           3      4    0
5         2        11           0           7           4      4    0
6         2         0           5           0           4      4    1

      native.country  income
1                38      0
3                38      0
4                38      0
5                38      0
6                38      0
```

```
[ ]: # Next, Concatenate df_categorical dataframe with original df (dataframe)

# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

```
[ ]:  age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  \
1   82  132870           9           0          4356           18
3   54  140359           4           0          3900           40
4   41  264663          10           0          3900           40
5   34  216864           9           0          3770           45
6   38  150601           6           0          3770           40

      workclass  education  marital.status  occupation  relationship  race  sex  \
1         2         11           6           3           1      4    0
3         2         5           0           6           4      4    0
4         2        15           5           9           3      4    0
5         2        11           0           7           4      4    0
6         2         0           5           0           4      4    1

      native.country  income
1                38      0
3                38      0
4                38      0
5                38      0
6                38      0
```

```
[ ]: # look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
```

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	age	30162 non-null	int64
1	fnlwgt	30162 non-null	int64
2	education.num	30162 non-null	int64
3	capital.gain	30162 non-null	int64
4	capital.loss	30162 non-null	int64
5	hours.per.week	30162 non-null	int64
6	workclass	30162 non-null	int64
7	education	30162 non-null	int64
8	marital.status	30162 non-null	int64
9	occupation	30162 non-null	int64
10	relationship	30162 non-null	int64
11	race	30162 non-null	int64
12	sex	30162 non-null	int64
13	native.country	30162 non-null	int64
14	income	30162 non-null	int64

dtypes: int64(15)

memory usage: 3.7 MB

```
[ ]: #convert target variable income to categorical  
df['income'] = df['income'].astype('category')
```

```
[ ]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

Int64Index: 30162 entries, 1 to 32560

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	age	30162 non-null	int64
1	fnlwgt	30162 non-null	int64
2	education.num	30162 non-null	int64
3	capital.gain	30162 non-null	int64
4	capital.loss	30162 non-null	int64
5	hours.per.week	30162 non-null	int64
6	workclass	30162 non-null	int64
7	education	30162 non-null	int64
8	marital.status	30162 non-null	int64
9	occupation	30162 non-null	int64
10	relationship	30162 non-null	int64
11	race	30162 non-null	int64
12	sex	30162 non-null	int64
13	native.country	30162 non-null	int64
14	income	30162 non-null	category

dtypes: category(1), int64(14)

memory usage: 3.5 MB

```
[ ]: # Importing train_test_split
from sklearn.model_selection import train_test_split
```

```
[ ]: # Putting independent variables/features to X
X = df.drop('income',axis=1)

# Putting response/dependent variable/feature to y
y = df['income']
```

```
[ ]: X.head(3)
```

```
[ ]:      age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  \
1    82  132870           9           0         4356           18
3    54  140359           4           0         3900           40
4    41  264663          10           0         3900           40

      workclass  education  marital.status  occupation  relationship  race  sex  \
1             2          11              6           3              1    4    0
3             2           5              0           6              4    4    0
4             2          15              5           9              3    4    0

      native.country
1                 38
3                 38
4                 38
```

```
[ ]: y.head(3)
```

```
[ ]: 1    0
3    0
4    0
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
[ ]: # Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪30,random_state=99)

X_train.head()
```

```
[ ]:      age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  \
24351  42  289636           9           0           0           46
15626  37  52465           9           0           0           40
4347   38  125933          14           0           0           40
23972  44  183829          13           0           0           38
26843  35  198841          11           0           0           35
```

	workclass	education	marital.status	occupation	relationship	race	\
24351	2	11	2	13	0	4	
15626	1	11	4	7	1	4	
4347	0	12	2	9	0	4	
23972	5	9	4	0	1	4	
26843	2	8	0	12	3	4	

	sex	native.country
24351	1	38
15626	1	38
4347	1	19
23972	0	38
26843	1	38

```
[ ]: # Importing decision tree classifier from sklearn library
from sklearn.tree import DecisionTreeClassifier

# Fitting the decision tree with default hyperparameters, apart from
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)
```

```
[ ]: DecisionTreeClassifier(max_depth=5)
```

```
[ ]: # Let's check the evaluation metrics of our default model

# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import
    classification_report,confusion_matrix,accuracy_score

# making predictions
y_pred_default = dt_default.predict(X_test)

# Printing classifier report after prediction
print(classification_report(y_test,y_pred_default))
```

	precision	recall	f1-score	support
0	0.86	0.95	0.91	6867
1	0.78	0.52	0.63	2182
accuracy			0.85	9049
macro avg	0.82	0.74	0.77	9049
weighted avg	0.84	0.85	0.84	9049

```
[ ]: # Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))
```

```
[[6553  314]
 [1039 1143]]
0.8504807161012267
```

```
[ ]: !pip install pydotplus
```

```
Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-
packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in
/usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)
```

```
[ ]: # Importing required packages for visualization
from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz


# Putting features
features = list(df.columns[1:])
features
```

```
[ ]: ['fnlwgt',
      'education.num',
      'capital.gain',
      'capital.loss',
      'hours.per.week',
      'workclass',
      'education',
      'marital.status',
      'occupation',
      'relationship',
      'race',
      'sex',
      'native.country',
      'income']
```

```
[ ]: # plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(dt_default, out_file=dot_data,
                feature_names=features, filled=True,rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

[]:



```
[ ]: # GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
```

```
# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'max_depth': range(1, 40)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=100),
                 param_grid={'max_depth': range(1, 40)}, scoring='accuracy')
```

```
[ ]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

```
[ ]:    mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      0.035181    0.011561      0.010106      0.002992
1      0.061280    0.018088      0.013943      0.008235
2      0.074714    0.026539      0.014324      0.013822
3      0.074829    0.025231      0.011173      0.008233
4      0.055147    0.004359      0.007180      0.002890

    param_max_depth  params  split0_test_score  split1_test_score  \
0                  1  {'max_depth': 1}      0.747810      0.747810
```


1	2	{'max_depth': 2}	0.812219	0.818612
2	3	{'max_depth': 3}	0.828558	0.834241
3	4	{'max_depth': 4}	0.832583	0.840871
4	5	{'max_depth': 5}	0.834241	0.844897

	split2_test_score	split3_test_score	split4_test_score	mean_test_score \
0	0.747573	0.747750	0.747750	0.747738
1	0.820507	0.825675	0.822833	0.819969
2	0.834478	0.836570	0.837518	0.834273
3	0.842529	0.842729	0.842255	0.840193
4	0.847265	0.842729	0.847466	0.843319

	std_test_score	rank_test_score
0	0.000087	39
1	0.004538	16
2	0.003115	12
3	0.003860	9
4	0.004858	7

```
[ ]: """
# plotting accuracies with max_depth
plt.figure()
plt.plot(scores["param_max_depth"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_max_depth"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("max_depth")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
"""
```

```
[ ]: '\n# plotting accuracies with
max_depth\nplt.figure()\nplt.plot(scores["param_max_depth"], \n
scores["mean_train_score"], \n          label="training
accuracy")\nplt.plot(scores["param_max_depth"], \n
scores["mean_test_score"], \n          label="test accuracy")\nplt.xlabel("max_de
pth")\nplt.ylabel("Accuracy")\nplt.legend()\nplt.show()\n'
```

```
[ ]: # GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

# specify number of folds for k-fold CV
```

```

n_folds = 5

# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)

```

```

[ ]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=100),
                param_grid={'min_samples_leaf': range(5, 200, 20)},
                scoring='accuracy')

```

```

[ ]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()

```

```

[ ]:
  mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      0.083994      0.003596      0.004026      0.000177
1      0.070918      0.003142      0.003753      0.000242
2      0.063995      0.001806      0.003521      0.000141
3      0.061234      0.004158      0.004317      0.000716
4      0.063998      0.009761      0.004074      0.000192

  param_min_samples_leaf  params  split0_test_score  \
0              5  {'min_samples_leaf': 5}      0.825716
1             25  {'min_samples_leaf': 25}      0.841819
2             45  {'min_samples_leaf': 45}      0.843003
3             65  {'min_samples_leaf': 65}      0.841108
4             85  {'min_samples_leaf': 85}      0.838030

  split1_test_score  split2_test_score  split3_test_score  split4_test_score  \
0      0.827848      0.819560      0.826149      0.818806
1      0.851291      0.839451      0.842018      0.849360
2      0.849159      0.846555      0.851018      0.851729
3      0.852711      0.845371      0.851492      0.838465
4      0.849159      0.845371      0.851492      0.842018

  mean_test_score  std_test_score  rank_test_score
0      0.823616      0.003696      10
1      0.844788      0.004651      6

```

2	0.848293	0.003194	1
3	0.845830	0.005589	2
4	0.845214	0.004834	3

```
[ ]: """
# plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_leaf"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_min_samples_leaf"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("min_samples_leaf")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
"""
```

```
[ ]: '\n# plotting accuracies with
min_samples_leaf\nplt.figure()\nplt.plot(scores["param_min_samples_leaf"], \n
scores["mean_train_score"], \n          label="training
accuracy")\nplt.plot(scores["param_min_samples_leaf"], \n
scores["mean_test_score"], \n          label="test accuracy")\nplt.xlabel("min_sa
mples_leaf")\nplt.ylabel("Accuracy")\nplt.legend()\nplt.show()\n'
```

```
[ ]: # GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                   cv=n_folds,
                   scoring="accuracy")
tree.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=100),
                  param_grid={'min_samples_split': range(5, 200, 20)},
                  scoring='accuracy')
```

```
[ ]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

```
[ ]:      mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      0.190965      0.053279      0.007489      0.002150
1      0.140156      0.022450      0.006022      0.000463
2      0.129533      0.010133      0.006719      0.000084
3      0.090867      0.022550      0.004043      0.000819
4      0.079408      0.004440      0.003670      0.000101
```

```
      param_min_samples_split      params  split0_test_score  \
0              5  {'min_samples_split': 5}      0.811982
1             25  {'min_samples_split': 25}      0.825006
2             45  {'min_samples_split': 45}      0.835188
3             65  {'min_samples_split': 65}      0.839451
4             85  {'min_samples_split': 85}      0.846081
```

```
      split1_test_score  split2_test_score  split3_test_score  split4_test_score  \
0      0.811035      0.818376      0.811701      0.808385
1      0.825243      0.830215      0.822596      0.827570
2      0.839687      0.830215      0.827333      0.838702
3      0.845844      0.837556      0.833728      0.843913
4      0.853895      0.838977      0.837281      0.845334
```

```
      mean_test_score  std_test_score  rank_test_score
0      0.812296      0.003296      10
1      0.826126      0.002581      9
2      0.834225      0.004783      8
3      0.840098      0.004360      7
4      0.844314      0.005898      6
```

```
[ ]: """
# plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_split"],
          scores["mean_train_score"],
          label="training accuracy")
plt.plot(scores["param_min_samples_split"],
          scores["mean_test_score"],
          label="test accuracy")
plt.xlabel("min_samples_split")
plt.ylabel("Accuracy")
```

```
plt.legend()
plt.show()
"""
```

```
[ ]: '\n# plotting accuracies with
min_samples_leaf\nplt.figure()\nplt.plot(scores["param_min_samples_split"], \n
scores["mean_train_score"], \n          label="training
accuracy")\nplt.plot(scores["param_min_samples_split"], \n
scores["mean_test_score"], \n          label="test accuracy")\nplt.xlabel("min_sa
mples_split")\nplt.ylabel("Accuracy")\nplt.legend()\nplt.show()\n'
```

```
[ ]: # Create the parameter grid
param_grid = {
    'max_depth': range(5, 15, 5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
    'criterion': ["entropy", "gini"]
}

n_folds = 5

# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
                           cv = n_folds, verbose = 1)

# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
[ ]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                param_grid={'criterion': ['entropy', 'gini'],
                            'max_depth': range(5, 15, 5),
                            'min_samples_leaf': range(50, 150, 50),
                            'min_samples_split': range(50, 150, 50)},
                verbose=1)
```

```
[ ]: # cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

```
[ ]:      mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      0.038696      0.001708      0.003295      0.000058
1      0.038543      0.001311      0.003356      0.000141
2      0.037763      0.001160      0.003247      0.000045
3      0.038969      0.002066      0.003636      0.000621
```

4	0.060942	0.002272	0.003482	0.000060
5	0.061919	0.001492	0.003888	0.000460
6	0.058421	0.003067	0.003427	0.000110
7	0.057515	0.001752	0.003373	0.000031
8	0.034163	0.000422	0.003496	0.000234
9	0.034110	0.000630	0.003269	0.000118
10	0.034014	0.000566	0.003180	0.000066
11	0.035852	0.002114	0.003935	0.000863
12	0.056857	0.002691	0.004146	0.000793
13	0.055483	0.000874	0.003322	0.000032
14	0.054844	0.005373	0.003306	0.000030
15	0.054681	0.003145	0.003366	0.000040

	param_criterion	param_max_depth	param_min_samples_leaf	\
0	entropy	5	50	
1	entropy	5	50	
2	entropy	5	100	
3	entropy	5	100	
4	entropy	10	50	
5	entropy	10	50	
6	entropy	10	100	
7	entropy	10	100	
8	gini	5	50	
9	gini	5	50	
10	gini	5	100	
11	gini	5	100	
12	gini	10	50	
13	gini	10	50	
14	gini	10	100	
15	gini	10	100	

	param_min_samples_split		params	\
0	50	{'criterion': 'entropy', 'max_depth': 5, 'min_...		
1	100	{'criterion': 'entropy', 'max_depth': 5, 'min_...		
2	50	{'criterion': 'entropy', 'max_depth': 5, 'min_...		
3	100	{'criterion': 'entropy', 'max_depth': 5, 'min_...		
4	50	{'criterion': 'entropy', 'max_depth': 10, 'min_...		
5	100	{'criterion': 'entropy', 'max_depth': 10, 'min_...		
6	50	{'criterion': 'entropy', 'max_depth': 10, 'min_...		
7	100	{'criterion': 'entropy', 'max_depth': 10, 'min_...		
8	50	{'criterion': 'gini', 'max_depth': 5, 'min_sam...		
9	100	{'criterion': 'gini', 'max_depth': 5, 'min_sam...		
10	50	{'criterion': 'gini', 'max_depth': 5, 'min_sam...		
11	100	{'criterion': 'gini', 'max_depth': 5, 'min_sam...		
12	50	{'criterion': 'gini', 'max_depth': 10, 'min_sa...		
13	100	{'criterion': 'gini', 'max_depth': 10, 'min_sa...		
14	50	{'criterion': 'gini', 'max_depth': 10, 'min_sa...		

15 100 {'criterion': 'gini', 'max_depth': 10, 'min_sa...

	split0_test_score	split1_test_score	split2_test_score \
0	0.834241	0.843950	0.840398
1	0.834241	0.843950	0.840398
2	0.834241	0.842529	0.840398
3	0.834241	0.842529	0.840398
4	0.842529	0.851527	0.847265
5	0.842529	0.851527	0.847265
6	0.845134	0.852475	0.847502
7	0.845134	0.852475	0.847502
8	0.834241	0.844897	0.847502
9	0.834241	0.844897	0.847502
10	0.834241	0.843476	0.844897
11	0.834241	0.843476	0.844897
12	0.843950	0.851291	0.849870
13	0.843950	0.851291	0.849870
14	0.836372	0.848449	0.843239
15	0.836372	0.848449	0.843239

	split3_test_score	split4_test_score	mean_test_score	std_test_score \
0	0.845097	0.845334	0.841804	0.004173
1	0.845097	0.845334	0.841804	0.004173
2	0.845097	0.845808	0.841615	0.004157
3	0.845097	0.845808	0.841615	0.004157
4	0.854334	0.853861	0.849903	0.004456
5	0.854334	0.853861	0.849903	0.004456
6	0.854098	0.845571	0.848956	0.003661
7	0.854098	0.845571	0.848956	0.003661
8	0.845097	0.847466	0.843841	0.004927
9	0.845097	0.847466	0.843841	0.004927
10	0.845097	0.845334	0.842609	0.004234
11	0.845097	0.845334	0.842609	0.004234
12	0.855045	0.855045	0.851040	0.004093
13	0.855045	0.855045	0.851040	0.004093
14	0.854098	0.846518	0.845735	0.005862
15	0.854098	0.846518	0.845735	0.005862

	rank_test_score
0	13
1	13
2	15
3	15
4	3
5	3
6	5
7	5

8	9
9	9
10	11
11	11
12	1
13	1
14	7
15	7

```
[ ]: # printing the optimal accuracy score and hyperparameters
print("best accuracy", grid_search.best_score_)
print(grid_search.best_estimator_)
```

best accuracy 0.8510400232064759

DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50)

```
[ ]: # model with optimal hyperparameters
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=10,
                                min_samples_leaf=50,
                                min_samples_split=50)

clf_gini.fit(X_train, y_train)
```


```
[ ]: DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50,
                           random_state=100)
```

```
[ ]: # accuracy score
clf_gini.score(X_test, y_test)
```

```
[ ]: 0.850922753895458
```

```
[ ]: # plotting the tree
dot_data = StringIO()
export_graphviz(clf_gini,
               out_file=dot_data, feature_names=features, filled=True, rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```




```
[ ]: # tree with max_depth = 3
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=3,
                                min_samples_leaf=50,
                                min_samples_split=50)

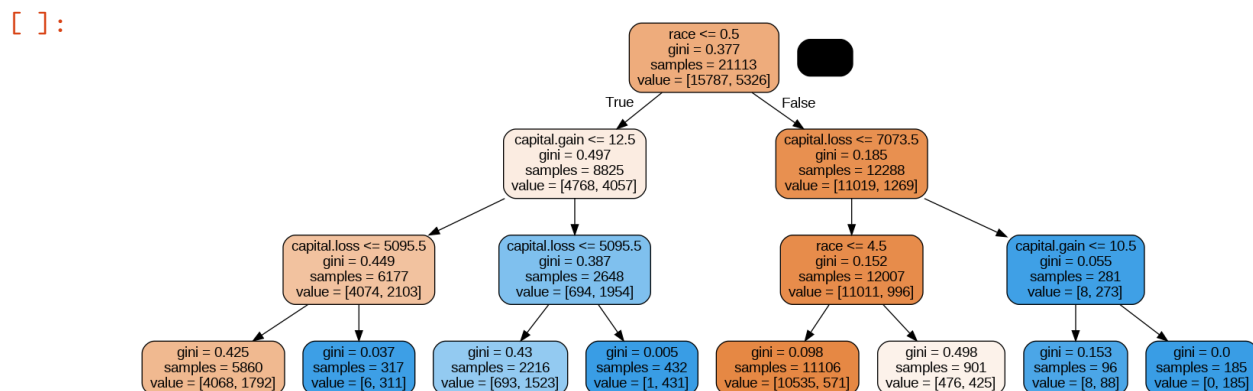
clf_gini.fit(X_train, y_train)

# score
print(clf_gini.score(X_test,y_test))
```

0.8393192617968837

```
[ ]: # plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(clf_gini,
               out_file=dot_data,feature_names=features,filled=True,rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



```
[ ]: # classification metrics
from sklearn.metrics import classification_report, confusion_matrix
y_pred = clf_gini.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	6867
1	0.77	0.47	0.59	2182
accuracy			0.84	9049
macro avg	0.81	0.71	0.74	9049
weighted avg	0.83	0.84	0.82	9049

```
[ ]: # confusion matrix  
      print(confusion_matrix(y_test,y_pred))
```

```
[[6564  303]  
 [1151 1031]]
```



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Conclusion:

1. Through the utilization of the Label Encoder technique, categorical attributes have been transformed from their original textual forms into numerical representations, assigning a unique integer to each distinct categorical value within each respective column. This transformation facilitates the integration of these categorical variables into machine learning algorithms, which typically require numerical inputs.
2. Hyperparameter tuning was conducted based on the decision tree model obtained, with a focus on key hyperparameters:
 - Max Depth: This parameter constrains the depth of the decision tree, guarding against excessive complexity and overfitting to the training data.
 - Min Samples Split: Setting a minimum number of samples required in a node before further splitting helps prevent overly specific decisions based on a small number of instances.
 - Min Samples Leaf: This parameter determines the minimum number of samples to be present in a leaf node, serving to discourage nodes with very few instances.
 - Criterion: The choice of criterion (e.g., "Gini impurity" or "entropy") defines the measure of split quality.
3. The model achieved an accuracy of approximately 84%, implying that it correctly predicted the class labels for 84% of instances in the test dataset.
4. Examining the confusion matrix, we find the following results:
 - True Positive (TP): 1031
 - True Negative (TN): 6564
 - False Positive (FP): 303
 - False Negative (FN): 1151

The confusion matrix reveals that the model performs reasonably well in predicting class 0, evident by the high count of true negatives and true positives. However, it faces challenges when predicting class 1.

5. Precision: Approximately 0.77, indicating that when the model predicts class 1, it is often correct. Roughly 77% of instances predicted as class 1 are genuinely class 1.
6. Recall: Approximately 0.47, implying that the model misses a notable number of actual class 1 instances. It correctly identifies only about 47% of all real instances belonging to class 1.
7. The F1 score for class 1 strikes a balance between precision and recall, offering a comprehensive view of the model's performance.